

# (Why We Need) Large Code Base Change Ripple Management in C++

Niall Douglas

# Contents:

1. What I am pitching: a new Boost library, and a possible motivating vision of a long-term future for C++ and Boost
2. Many (contentious) claims as to why ... 😊

**What I am pitching: a low-level embedded graph database for Boost**

# The proposed embedded graph database

- All first tier content is a standard file which can be opened, mmaped etc
- Takes advantage of filing system specific features such as extents, metadata/data journaling, hole punching, copy on write, bitrot self healing etc
- Strong versioning and MVCC concurrency
- Per-graph content protection (e.g. parity healing)
- Content addressable with a per-graph hash of your choice

# The proposed embedded graph database

- Per-graph optional ACID transactions
- Per-graph arbitrary indexers (e.g. Boost.Graph, SQLite3, ZIP etc)
- Network shardable to other copies with interrupted partial copy resumption
- Objects can be executable (in fact is self hosting)
- Uses an algorithm very close to git
- Designed to act like increasing mount points of 'database-ness' overlaid onto the filesystem

# The proposed embedded graph database

- Performance is expected to be within two and five orders of magnitude slower than the big iron graphstores
- Write transaction performance shouldn't be lower than ten per second hopefully
- All designed to work during very early process bootstrap i.e. before shared libraries are loaded
- There is nothing close in existing software - this design and its abilities are very unique

It is really more of a  
“generic data persistence  
library”



# Want more detail?

There is a 25,000 word accompanying position white paper on ArXiv at <http://arxiv.org/abs/1405.3323>:

- What makes code changes ripple differently in C++ to other languages?
- Why hasn't C++ replaced C in most newly written code?
- What does this mean for C++ a decade from now?
- What C++ 17 is doing about complexity management
- What C++ 17 is leaving well alone until later: Type Export
- Detail about the embedded graph database design
- Two example killer applications for such a graph database namely:
  - a. An example C++ object components design (similar to Bandela's)
  - b. An example extensible Filesystem design



# What does any of this have to do with the price of fish?

- What does this have to do with change ripple management?
- Or Boost?
- Or C++?
- Or anything?

# Let the contention begin!

I am now going to articulate a (motivating) vision for a long term future goal of C++ and Boost which explains why we might absolutely need one of these databases soon

I will then make a series of supporting claims most of which will be contentious (and hence I place them last!)

**What is this (motivating)  
vision of the future of C++  
and Boost?**

# A post-C++ 17 goal:

I'd like to see a world where we can write C++ as if everything in the solution (including Python, Lua, PHP etc, including C++ in closely related processes) is header only, no matter the size of the program

(It is a natural result of a complete ABI management solution)

# Example:

```
class Foo { virtual void boo(int a); };
```

A rule in the graph database says that when this type is changed, it should be reflected via `std::reflect` into a Python binding:

```
class Foo:  
    @accepts(int)  
    @returns(None)  
    def boo(self, a)
```

# Example:

Let's break Foo's ABI:

```
class Foo { virtual void boo (double); };
```

When you hit compile in C++ for each use of `Foo::boo()` in Python code you get:

```
TypeWarning: 'boo' method accepts (float), but was given  
(int)
```

# To clarify:

- C++ moves from a source file compilation model to a type graph compilation model (similar to exported templates)
- The type graphs are compiled a bit like GLSL shaders into many tiny C++ Modules i.e. bits of precompile all put into the graphstore
- Reflection (runtime) equals a graph query
- To bootstrap a C++ process equals visiting a graph query for all the matching tiny C++ Modules



# Consequences:

- Instant notification of breakages from a code change no matter how far away
- Optimally minimal rebuild (and can dispense with external build tools)
- Optimal optimisation which can be pushed onto a batch pass/cloud compute

# Consequences:

- Easy components via ripple propagation rules in the graphstore
- No longer pollutes all over the C symbol table
- Finally a real ABI management solution
- Other programming languages would be very interested in this

Why???

# Claims (to which I shall return):



1. C++ is in relative decline
2. Boost is in both absolute and relative decline

Therefore:


3. We need to return to becoming a better systems programming language
4. We need “signature projects” for C++ 11/14

Why this instead of  
including more code per-  
compiland?

# Where hardware is going soon



# Where hardware is going soon





# My best speculations on effects:

I therefore claim these likely outcomes in the future:

1. The cost of including ever more source code per-compilant stops being sunk by transistor density growth
2. Therefore build times start to rise and stop falling with time
3. Therefore C++ starts to look more like in the 1990s with small compiles and large links
  - Except with all our fancy modern C++ techniques

# My best speculations on effects:

Also:

4. There will be a return to growth for systems programming languages as software is refactored to cope with linear growth CPU and RAM but still exponential growth storage
5. I claim this will happen around 2017-2020 if present trends continue and no surprises turn up  
(Mass production of Graphene or Phosphorine transistors won't be ready by 2020 at present rates of R&D)

Will C++ be that majority  
choice of systems  
language?

# The present structural revolution

## Compare Languages

Monthly Commits

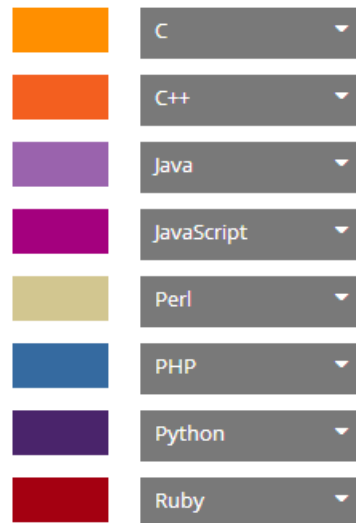

Monthly Contributors

Monthly Lines of Code Changed

Monthly Projects

### Monthly Contributors

The lines show the number of developers who have contributed at least one line of code in each month. [More](#)



**Claim:**  
**C++ has been in relative  
decline for a decade**  
(with a pause 2009-2011)

# The present structural revolution

## Compare Languages

Monthly Commits


Monthly Contributors

Monthly Lines of Code Changed

Monthly Projects

### Monthly Contributors (Percent of Total)

The lines show the number of developers who have contributed at least one line of code in each month. [More](#)



# The present structural revolution

## Compare Languages

Monthly Commits

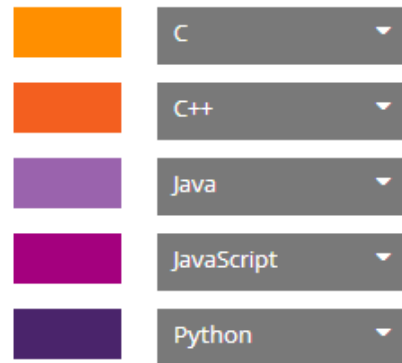

Monthly Contributors

Monthly Lines of Code Changed

Monthly Projects

### Monthly Projects (Percent of Total)

The lines show the count of projects with at least one line of code changed in a month. [More](#)





# Possible reasons why C++ is in relative decline

And why should we care?

# Why C++ is in relative decline?

1. C++ is no longer a general purpose programming language - it's a *niche specialist* language suited for:
  - a. Low latency (async etc)
  - b. Maximum performance (maths etc)
  - c. Gluing application and service code written in other languages like Python or C# together

Note that C is good at all of the above too, and still remains more popular in open source than C++ for new code

# Why C++ is in relative decline?

2. C++ has stopped trying to be the best systems programming language possible
  - C++ 11/14 adds a ton of great stuff BUT ...
    - Did any of it persuade someone like Linus that C++ might be tolerable in the Linux kernel?
    - Do the Python/Ruby/Lua/PHP interpreter guys look at C++ 11/14 and go “wow that transforms our use case for C++ over C”?

# Why should we care?

If:

- a. C++ remains best in class for high performance math
- b. C++ remains very strong in low latency async
- c. BUT C remains preferred to C++ as a systems programming language

# Why should we care?

Then, assuming the previous claims are true, a reasonable prediction of the future is:

1. C (or some extension thereof) gets the majority of post-exponential hardware systems language growth
2. C++ becomes ever more like Haskell, with only a rarified programmer elite able to touch it


Is that what we want?

**Claim:**  
**Since 2011 Boost is in**  
**both an absolute &**  
**relative decline**

# The Decline of Boost

Commits per Month


Zoom 1 yr 3 yr 5 yr 10yr All






# The Decline of Boost


Number of Contributors



Libraries added to Boost



# The Decline of Boost



# Possible reasons why Boost is in decline

And why should we care?

# Why Boost is in decline?

1. The C++ 11/14 standard library now can do what people used to need Boost for
  - Unfixed bugs in Boost means that people simply switch to C++ 11/14 instead if they can
  - People think a particular Boost library needs all of Boost as a dependency - a real showstopper
  - Boost is seen as simply no longer relevant

# Why Boost is in decline?

To quote a highly respected engineer from this very conference who said a few days ago:

“Boost used to be about all the stuff you really wanted in the standard. Now Boost looks like all the stuff that wasn’t good enough to get into the standard”

- somebody well known (not me)

# Why Boost is in decline?

2. Boost has become two mutually incompatible sets of libraries:
  - a. The C++ 11 STL emulation library for C++ 98
  - b. A set of libraries which push the boundaries of C++, as Boost once used to in the 1990s
    - With the latter being suffocated of late

# Why Boost is in decline?

3. Most of the interesting C++ 11 libraries on the internet appear to have no interest in joining Boost (with a few honourable exceptions)

I personally find that very scary

# Why Boost is in decline?

4. Boost makes no attempt to preserve ABI stability, and therefore is not welcome in large stable code bases
  - None of the improvements in C++ 11/14 do anything for change ripple management, so even mild ABI breakage is intolerable and therefore Boost is banned/pinned to some ancient version



# Why should we care?

Simple answer:

How many of the changes to C++ 11/14 standard over C++ 98/03 originated in Boost?

Assertion:

C++ ought to return to  
trying to become a better  
systems programming  
language

**Assertion:**  
**We want Boost to**  
**continue to lead out the**  
**future of C++**  
**(and therefore all systems programming)**

# So what?

What if one or all of the earlier assertions is false?

What if none of the issues described is a real problem?

What if all this is merely hand wavy nonsense?

An embedded graph database is still extremely useful:

- In 2014 it is still too hard for more than one process to write to many files concurrently
- In 2014 it is still too easy to lose data

# Want more detail?

There is a 25,000 word accompanying position white paper on ArXiv at <http://arxiv.org/abs/1405.3323>:

- What makes code changes ripple differently in C++ to other languages?
- Why hasn't C++ replaced C in most newly written code?
- What does this mean for C++ a decade from now?
- What C++ 17 is doing about complexity management
- What C++ 17 is leaving well alone until later: Type Export
- Detail about the embedded graph database design
- Two example killer applications for such a graph database namely:
  - a. An example C++ object components design (similar to Bandela's)
  - b. An example extensible Filesystem design