

Iterators Will Stay

A Survey of Range Libraries

Sebastian Redl
sebastian.redl@getdesigned.at

U had a LOLcat...

but I ate it!

Range Library

```
vector<int> v = { 4, 2, 5, 2, 5, 2, 1, 3 };
```

```
sort(v);
```

```
copy(equal_range(v, 2), insert<int>(cout));
```


Adapters

- Wrappers for iterators (or ranges)
- Modify primitive operations
 - Skip elements on increment (filtering, striding)
 - Apply function on dereference (projection)
 - etc.

SG 9

- Established late 2012 / early 2013
- Headed by Marshall Clow
- “The goal of this study group (SG9) is to research the idea of adding ranges to a future version of the C++ standard library, and to create a proposal for the committee to consider.”

Goals

- Convenience
- Efficiency
- Safety
- Reuse of old code

Reuse of Old Code

- Use standard iterators with new algorithms

Reuse of Old Code

- Use standard iterators with new algorithms
- Use new ranges with old algorithms

Reuse of Old Code

- Use standard iterators with new algorithms
- Use new ranges with old algorithms
- Massively inhibits freedom of design

Reuse of Old Code

- Use standard iterators with new algorithms
- Use new ranges with old algorithms
- Massively inhibits freedom of design
- Probably a hard requirements for SG 9 result

Convenience

- Primary motivation for ranges

```
sort(v.begin(), v.end());  
// vs.  
sort(v);
```

Convenience

- Primary motivation for ranges

```
sort(v.begin(), v.end());  
// vs.  
sort(v);
```

```
copy(equal_range(v, 2), insert<int>(cout));  
// vs.  
auto r = equal_range(v.begin(), v.end(), 2);  
copy(r.begin(), r.end(),  
      ostream_iterator<int>(cout));
```

Convenience

- Iterators are hard to write
 - Comparison is often unintuitive

Convenience

- Iterators are hard to write
 - Comparison is often unintuitive

```
bool equal(iterator that) const
{
    return str_
        ? (that.str_ == str_ ||
           (!that.str_ && !*str_))
        : (!that.str_ || !*that.str_);
}
```

Convenience

- Iterators are hard to write
 - Comparison is often unintuitive
- Adapters are hard to use
 - Every iterator needs to be wrapped

Convenience

- Iterators are hard to write
 - Comparison is often unintuitive
- Adapters are hard to use
 - Every iterator needs to be wrapped

```
copy(make_move_iterator(  
    make_reverse_iterator(first)),  
    make_move_iterator(  
        make_reverse_iterator(last)),  
    target);  
copy(make_move_range(make_reverse_range(r)),  
    target);
```

Convenience

- Iterators are hard to write
 - Comparison is often unintuitive
- Adapters are hard to use
 - Every iterator needs to be wrapped
- Adapters are hard to write
 - Avoiding undefined behavior

Efficiency

- Ranges must be as efficient as iterators
- Minimal overhead over hand-written algorithm

Efficiency

- Iterators duplicate information
 - Projection iterators: function in every iterator
 - Filter iterators: must know end iterator

Efficiency

- Iterators duplicate information
 - Projection iterators: function in every iterator
 - Filter iterators: must know end iterator
- Ranges require good calling convention
 - Iterators may fit into register, ranges don't
 - Needs compiler to split struct across registers

Safety

- Detect out-of-bounds access
- Avoid invalidation traps

Sample Libraries

- Boost.Range
- Eric Niebler's Range v3 (Iterables)
- Phobos std.range (D standard library)
- libaccent

The Big Divide

- | | |
|---------------------------------------|---------------|
| • Boost.Range | Use Iterators |
| • Eric Niebler's Range v3 (Iterables) | |
- | | |
|---|--------------|
| • Phobos std.range (D standard library) | |
| • libaccent | No Iterators |

Boost.Range

- Range is anything with begin/end
- Iterator is still main primitive

Boost.Range

- Range is anything with begin/end
- Iterator is still main primitive
- Perfect compatibility

Boost.Range

- Range is anything with begin/end
- Iterator is still main primitive
- Perfect compatibility
- Inherits all efficiency and safety downsides
- Inherits some convenience downsides

Boost.Range

- Range is anything with begin/end

```
vector<int> v = { 4, 2, 5, 2, 5, 2, 1, 3 };  
rg::sort(v);
```

- And it doesn't have anything else

Boost.Range

- Iterator is still main primitive

```
template <typename Range>
void sort(Range& r) {
    std::sort(boost::begin(r),
              boost::end(r));
}
```

Boost.Range

- Inherits all efficiency and safety downsides
 - No protection beyond what iterators have
 - Adapted at iterator level: space explosion
 - Passing ranges by value would mean big objects
 - Ranges are not meant to be passed by value

Boost.Range

- Inherits some convenience downsides
 - Range adapters just produce iterator adapters
 - Still need to write iterator adapters

Eric Niebler's Range v3

- Iterator is still main primitive
- Iterator pairs need not be same type

Eric Niebler's Range v3

- Iterator is still main primitive
- Iterator pairs need not be same type
- Iterable is anything with begin/end
- Range is a homogenous Iterable

Eric Niebler's Range v3

- Iterator pairs need not be same type
 - `istream_iterator` knows when it's done
 - An “end” `istream_iterator` is ugly to implement
 - Instead, have a sentinel “end” iterator
 - Comparison with sentinel is real iterator's `is_done`

Eric Niebler's Range v3

- Iterator pairs need not be same type
 - Can represent arbitrary end predicate
 - Only works up to forward ranges
 - Bidirectional ranges must be homogenous
 - Can this represent counted ranges effectively?

Lifetime issues

- Is this code valid?

```
auto foo(const vector<int>& v) {  
    auto r = reverse(  
        transform(SomeFunction(),  
            v));  
    auto it = find(r, 99);  
    if (it == r.end()) return it;  
    return next(it);  
}
```

In Boost.Range? In Range v3?

Lifetime issues

- Boost.Range
 - Don't know. Doesn't document it.
- Range v3
 - No. Iterators depend on ranges.

The Big Divide

- Iterator pairs can be very awkward
- Relaxing requirements solves some problems
- Andrei Alexandrescu: Iterators Must Go!

Phobos std.range

- Implements ideas of Alexei Alexandrescu's "Iterators Must Go" keynote of 2009
- Range is the main primitive

Phobos std.range

- Implements ideas of Alexei Alexandrescu's "Iterators Must Go" keynote of 2009
- Range is the main primitive
- Strictly less powerful than bidirectional iterators
 - Ranges can never grow

Phobos std.range

- Implements ideas of Alexei Alexandrescu's "Iterators Must Go" keynote of 2009
- Range is the main primitive
- Strictly less powerful than bidirectional iterators
 - Ranges can never grow
- Sufficient for all purposes, but can be unintuitive
 - What does find() return?

Range Operations

- Forward traversal
 - empty?
 - access first element
 - drop first element

Range Operations

- Forward traversal
 - empty?
 - access first element
 - drop first element
- Bidirectional traversal
 - access last element
 - drop last element

Range Operations

- Forward traversal
 - empty?
 - access first element
 - drop first element
- Bidirectional traversal
 - access last element
 - drop last element
- Random access traversal
 - Drop arbitrary number of elements on either side

Phobos std.range

- What does find() return?

Phobos std.range

- What does find() return?
 - find() returns range from found element

Phobos std.range

- What does find() return?
 - find() returns range from found element
 - findSkip() returns range after found element

Phobos std.range

- What does find() return?
 - find() returns range from found element
 - findSkip() returns range after found element
 - findSplit() returns before match, match, after match

Phobos std.range

- What does find() return?
 - find() returns range from found element
 - findSkip() returns range after found element
 - findSplit() returns before match, match, after match
 - until() returns until before or after match (flag)

Phobos std.range

- What does find() return?
 - find() returns range from found element
 - findSkip() returns range after found element
 - findSplit() returns before match, match, after match
 - until() returns until before or after match (flag)
- Suddenly four (seven?) algorithms

libaccent

- Phobos-style Range
- Additional primitive to represent position

libaccent

- Phobos-style Range
- Additional primitive to represent position
 - Knows whether it refers to something

libaccent

- Phobos-style Range
- Additional primitive to represent position
 - Knows whether it refers to something
 - Can be dereferenced

libaccent

- Phobos-style Range
- Additional primitive to represent position
 - Knows whether it refers to something
 - Can be dereferenced
 - Can be used to cut ranges short

libaccent

- Phobos-style Range
- Additional primitive to represent position
 - Knows whether it refers to something
 - Can be dereferenced
 - Can be used to cut ranges short
 - Cannot be incremented

libaccent

- Phobos-style Range
- Additional primitive to represent position
 - Knows whether it refers to something
 - Can be dereferenced
 - Can be used to cut ranges short
 - Cannot be incremented

```
if (auto p = find(rng, is(42))) {  
    std::cout << *p << '\n';  
    auto before = until(rng, p);  
    auto after = after(rng, p);  
}
```


Lifetime Issues

- Ranges always used by value
 - Containers are not ranges

Lifetime Issues

- Ranges always used by value
 - Containers are not ranges
- Positions do not depend on ranges

Lifetime Issues

- Ranges always used by value
 - Containers are not ranges
- Positions do not depend on ranges
- Invalidation only happens if underlying sequence goes away
 - Might be garbage-collected class in D

Iterators vs No Iterators

- Ranges more flexible in implementation
 - Delimited, counted, infinite ...

Iterators vs No Iterators

- Ranges more flexible in implementation
 - Delimited, counted, infinite ...
- Iterators more flexible for algorithms
 - Two iterators == one range

Iterators vs No Iterators

- Ranges more flexible in implementation
 - Delimited, counted, infinite ...
- Iterators more flexible for algorithms
 - Two iterators == one range
 - Three iterators == three ranges
 - range + find() result
 - Four iterators == six ranges
 - range + equal_range() result

Iterators vs No Iterators

- Ranges are safer
 - Know when they are empty
 - Can never grow

Iterators vs No Iterators

- Ranges are safer
 - Know when they are empty
 - Can never grow
- Ranges work more easily with adapters
 - Range adapters easier than iterator adapters
 - No iterators with dependent lifetimes

Iterators vs No Iterators

- Iterators are harder to invalidate
 - List splice can invalidate ranges
 - List modification invalidates ranges that hold count

Iterators vs No Iterators

- Iterators are harder to invalidate
 - List splice can invalidate ranges
 - List modification invalidates ranges that hold count
- `any_range` far simpler than `any_iterator`

Iterators vs No Iterators

- Iterators are harder to invalidate
 - List splice can invalidate ranges
 - List modification invalidates ranges that hold count
- `any_range` far simpler than `any_iterator`
- Iterators are used by existing code

Iterator-Based Range Libraries

Library	Adapters	Invalidation	Copy Semantics
Boost.Range	Ranges very thin wrapper	Undocumented Iterators held by value	Reference semantics
Range v3 (Eric)	Ranges hold adapter logic	Iterators depend on ranges	Reference semantics
Ranges (Chandler)	Ranges own elements and hold logic	Iterators depend on ranges	Value semantics
Ranges (Arno Schödl)	???	Iterators independent of ranges	Non-copyable

Discussion

- Boost.Range:

http://www.boost.org/doc/libs/1_55_0/libs/range/doc/html/index.htm

- Range v3:

<http://ericniebler.com/2014/02/16/delimited-ranges/>

- Chandler's Ranges: ???

- Phobos:

http://dlang.org/phobos/std_range.html

- libaccent (use “rewrite” repository):

<https://code.google.com/p/libaccent/>