

# Ownership of Memory

## Guidelines for Dynamic Allocation

By David Stone

# The Basics

- Automatic storage ("the stack")

```
T t;
```

- Free store ("the heap")

```
new T();
```

```
std::make_unique<T>();
```

```
std::make_shared<T>();
```

# Hierarchy of Ownership

- 1) Automatic variables
- 2) `std::unique_ptr`
- 3) `std::shared_ptr`

Note: raw pointers are not on this list

# "Large" values and rvalues

- For this presentation, a "large" value refers to the value returned by `sizeof`
- `std::vector<T>` is small
  - 24 bytes on 64-bit
  - Regardless of number of elements
- `std::array<char, 1000>` is large
  - 1000 bytes

# "Large" values and rvalues

- Small values are cheap to move
- Large values are expensive to move

Two light blue squares are positioned vertically on the right side of the slide, one above the other.

# Why dynamic?

# Run-time sized collections

- `std::vector`
- `std::deque`
- `std::map`

# Polymorphism

- Dynamic allocation is required for runtime polymorphism in function returns
  - `auto create_object() -> std::unique_ptr<Base>;`
- Usually required for member and local variables
  - `std::unique_ptr<Base> m_some_object;`
- Generally not required for function parameters



# When move is not an optimization of copy

- Some objects must remain at the same address (reference stability)
  - `std::mutex`
  - Anything that other objects reference
  - Multithreaded code
- `std::unique_ptr<T>` is always movable, even if `T` is not

# When move is an optimization of move

- `std::unique_ptr<T>` is always fast to move, even when `T` is not
  - `std::array<int, 10000>` is slow to move
- `std::unique_ptr<T>` requires constant stack space of one pointer
  - `std::array<std::array<int, 1024>, 1024>` will probably overflow your stack

# Cache-friendliness, sequence

- `std::deque<T>`
- `std::list<T>`
- `std::vector<std::unique_ptr<T>>`
  - `moving_vector`
- `std::vector<T>`


# Source of data

- <https://bitbucket.org/davidstone/containers>
- Compiled with g++ 4.8.2
  - Ofast
  - march=native
  - fipa-pta
  - funsafe-loop-optimizations
  - flto

# Cache-friendliness, sequence

Time to search and insert into middle


1-byte elements



# Cache-friendliness, sequence

Time to search and insert into middle


10-byte elements



# Cache-friendliness, sequence

Time to search and insert into middle


40-byte elements



# Cache-friendliness, sequence

Time to search and insert into middle

50-byte elements






# Cache-friendliness, sequence

Time to search and insert into middle


60-byte elements



# Cache-friendliness, sequence

Time to search and insert into middle


100-byte elements



# Cache-friendliness, sequence

Time to search and insert into middle


200 byte elements



# Cache-friendliness, sequence

Time to search and insert into middle


500-byte elements



# Cache-friendliness, sequence


Time to search and insert into middle

1000 byte elements



# Cache-friendliness, sequence

1,000,000 elements





# Explanation of results

- Locality of reference
  - Explains plain vector's performance

# Explanation of results

- Pipelining





## std::list

- If you are iterating, do not use std::list
- If you are not iterating, do not use std::list

# Cache-friendliness, associative

- `std::map<T>`
- sorted `std::vector<T>`
  - `unstable_flat_map`
- sorted `std::vector<std::unique_ptr<T>>`
  - `stable_flat_map`

# Cache-friendliness, associative

No performance graphs.



=(

## Cache-friendliness, associative

- `unstable_flat_map` provides the fastest lookup
- `std::map` provides the fastest insertion
- For small element sizes, `unstable_flat_map` is faster for everything else (including batch insertion)
  - Small is around 100 bytes

# Cache-friendliness, associative

- `stable_flat_map` was never the best choice
- Why does it perform so much worse at associative tasks?



# Optional values

- What would a function look like that needs to output an optional value?
  - What does it look like if it just needs to return "by reference"
- Which interface to use?
  - Which is easier?
  - Which is faster?

# Optional values

- Typically, use something like the proposed `std::optional`
  - <https://github.com/akrzemi1/Optional>
- `std::optional<T>` is always at least as large as `T`
- Null `std::unique_ptr<T>` is `sizeof(T *)`

# Do not be afraid of "by value"

- The compiler will elide copies
  - Return value optimization
- Most moves are cheap

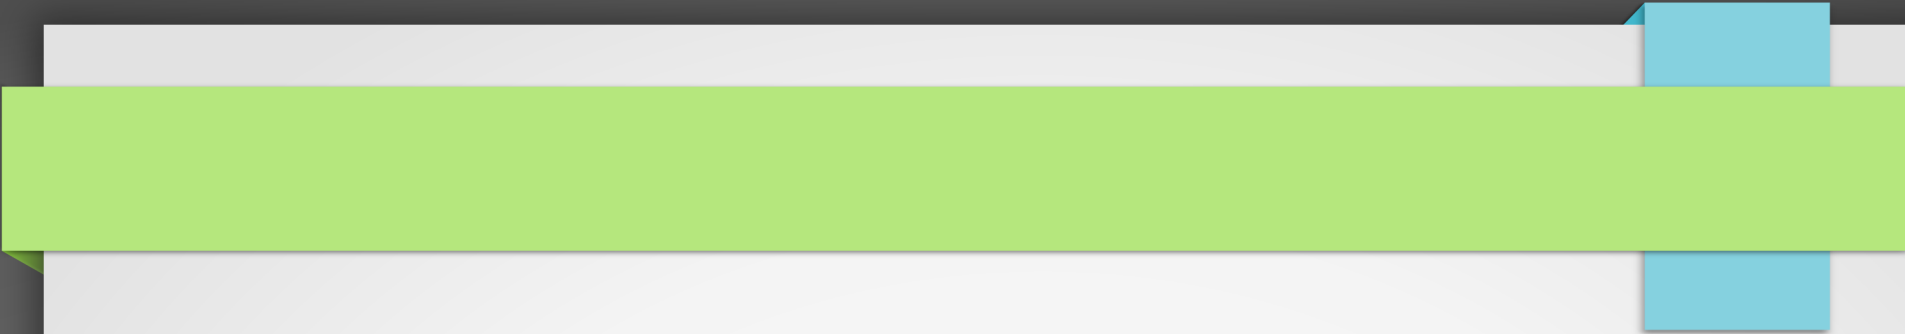


# Returning a dynamically allocated value

- `auto u() -> std::unique_ptr<T>;`
- `auto s() -> std::shared_ptr<T>;`

# Constructing a shared\_ptr

- `std::shared_ptr<T>(new T());`
- `std::make_shared<T>();`

- 
- The importance of `std::make_shared`
  - The difference between libraries and applications

# Value semantics

- `std::shared_ptr` is last on the hierarchy
- indirection is powerful
  - With great power comes great difficulty in comprehension
- `std::shared_ptr<T const>` is not as bad

# Scope-Bound Resource Management

- Destructor is a fundamental part of C++
- Structured code automatically creates nested life times
  - Functions, not goto everywhere
- Do not use `std::shared_ptr` to avoid structuring code
- `std::shared_ptr` is used with multi-threading
  - When you are modeling something that must hold onto data, but the duration is dependent on run-time factors

# Raw Pointers

- Raw pointers never own memory
- Smart pointers do not deprecate raw pointers
- Raw pointers reference memory
- A pointer is an optional reference

# Compilation firewall

```
#include <lots_of_headers>

class Class {
public:
    functions();
private:
    Thing1 m_1;
    Thing2 m_2;
    ...
    ThingN m_n;
};
```

# Compilation firewall

```
class Class {  
public:  
    functions();  
private:  
    class Impl;  
    std::unique_ptr<Impl>  
};
```



# Compilation firewall

- Minimize compile times
- Minimize recompilation
- Stable ABI
- If used with care, can work around binary compatibility problems in third-party libraries

# Summary

- Use a smart pointer only for:
  - Dynamic polymorphism when you cannot achieve this without dynamic allocation
  - Minimizing compilation dependencies (PImpl)
  - Optional values that could be very large
  - Enabling moves on non-movable types
  - Optimizing moves on slow-to-move types
- In all other cases, use an automatic variable
- Prefer cache-friendly data structures